YALLO TALENT

The corridor runs both ways: enterprise platform talent across the UK, Saudi Arabia and the UAE

Five platform families, three markets, one measurement. The finding that holds them together is not that one end of the corridor is short of people. It is that the two ends specialise, and in opposite directions.

SOURCE: LINKEDIN TALENT INSIGHTS · AUGUST 2026

? Ask Yallo Talent

The argument

The usual account of enterprise talent between the UK and the Gulf has one direction: the Gulf is short of specialists, so specialists come in. The measurement does not support that account, and the way it fails is more useful than the account itself.

Oracle is the counter-example that breaks the pattern open. 45.3% of the Oracle professionals measured sit in Saudi Arabia or the UAE, the highest Gulf share of any family here, and Oracle Financials is declared more widely in the Gulf than in the UK. Salesforce runs the other way: the UK pool is 4,326 against 65 in Saudi Arabia, and Marketing Cloud is far more widely declared in the UK.

So the corridor is not a shortage at one end and a surplus at the other. The two ends have specialised, and a programme's staffing problem depends on which platform it is running and where. An Oracle finance programme in Riyadh is in the deep end of its market. A Salesforce commerce programme in the same city is not, and no amount of local recruitment effort changes that.

This is a more defensible position than scarcity alone, and a harder one for a competitor to copy, because it requires knowing what a programme actually needs in each workstream as well as what the market declares. Two of the sharpest findings sit inside families that look healthy from the outside: Oracle payroll is the scarcest named skill in the entire extract and sits in the workstream with a legal date on it, and the cloud pool is so broadly declared that the constraint moves from finding people to telling them apart.

The practical conclusion is the same in every case and it is about timing rather than effort. The roles that will have to come from outside are knowable at business case. Naming them then costs a paragraph in a plan. Discovering them at week six costs the date.

How to read these numbers

Figures are drawn from a LinkedIn Talent Insights extract read on 2 August 2026, covering the United Kingdom, Saudi Arabia and the United Arab Emirates. Every figure carries that date.

Skills on LinkedIn are self-declared. A profile listing a platform is not a person who has delivered with it. This is the dataset's weakest property, and it is also the argument: the distance between a declared skill and a screened one is the work.

The named skill counts inside a family overlap, because one professional declares several. They are independent shares of the same pool and can sum past 100%. No chart here presents them as parts of a whole.

Baselines are not comparable between families. The size of a pool reflects how wide the title net was, not how large a market is, so no chart places two families on one axis.

This is supply, not demand. It counts professionals who exist, not roles that are open, and it supports no claim about how long a role takes to fill or how hard it is to hire.

The extract has no Gulf column. Where a figure compares the Gulf with the UK it is the unweighted mean of the Saudi Arabian and Emirati shares, stated so the construction can be checked.

The research covers the five platform and capability families the extract measures. It is not a complete view of the desks Yallo staffs, and nothing here should be read as coverage of the ones it does not name.

The specialisation asymmetry

Oracle carries 45.3% of its measured pool in the Gulf. Oracle Financials is declared by 48.5% of the Gulf pool against 40.1% in the UK.

Salesforce inverts it. Marketing Cloud is declared by 60.1% of the UK pool and 46.4% across the Gulf.

Two families, two directions, one dataset. The corridor exists because programmes need skills that are unevenly distributed, and the unevenness does not have a consistent sign.

SAP: the pool is weighted toward the phase that ends

Data migration is the largest declared skill in the regional SAP pool and security is the thinnest. A rollout can staff the workstream that finishes and will struggle to staff the one that never does.

Data migration is the largest named skill in the SAP pool: 27.4% of 7,922 professionals across the UK, Saudi Arabia and the UAE, 2,169 people. Security is the thinnest at 11.0%.

That ordering is not an accident of sampling. Migration is the workstream with a date on it, visible on every programme plan, and the skill a consultant can name after one cutover. Security is the workstream that runs after go-live and never finishes, and it is the line that gets cut at business-case stage when the number needs to come down.

The regional market has, in effect, declared its skills around go-live events. A rollout can therefore staff its migration workstream from the pool. It will struggle to staff its security workstream from the same pool, and the shortfall arrives at the point in the programme where there is least room to absorb it.

Migration is highest in the UAE at 32.1%, against 27.5% in the UK. Integration runs 4.3 percentage points higher in the Gulf than in the UK.

Integration running higher in the Gulf is consistent with a region that bought a great many platforms in a short period and then had to join them together. That is a coherent picture of a market's recent history rather than a curiosity, and it is useful: an integration-heavy programme has more to work with regionally than the headline pool size suggests.

Conclusion: The SAP market has declared its skills around go-live events, so a programme staffs its migration workstream from the pool and has to plan for its security workstream.

Oracle: the Gulf is the deep end, and payroll is still the problem

Oracle is the one family where the Gulf holds the deeper pool. Depth is not evenness: inside it, payroll is the scarcest named skill in the entire dataset, and payroll is the workstream that cannot slip.

45.3% of the Oracle professionals measured sit in Saudi Arabia or the UAE — the highest Gulf share of any family in the set. Financials runs 8.4 percentage points higher in the Gulf than in the UK, and Fusion 6.5 percentage points higher.

This matters because the corridor is usually described in one direction, as though the Gulf were short of everything and the UK had a surplus. On Oracle it is the other way round. The two ends of the corridor specialise; neither is simply the shallow end.

103 professionals across all three markets declare Oracle payroll: 46 in the UK, 29 in Saudi Arabia, 28 in the UAE. That is 2.0% of the Oracle pool, and the smallest named skill in the entire extract.

Payroll is also the workstream that cannot slip. It is legally dated, it touches every employee, and it is the one part of a Fusion programme where a missed date is visible outside the programme on the same day it happens.

So the planning consequence is specific rather than general. Finance can be staffed from the Gulf pool with confidence. Payroll cannot be staffed from any of these three markets on the assumption that a search will simply find someone, and a plan that assumes it will is carrying an unpriced risk in its most exposed workstream.

Treating payroll as a corridor role from day one means deciding, at business case, that the specialist will be engaged from outside the country and that someone has to be able to employ them there. That is entity and payroll friction, and it is precisely what an employer of record removes.

The connection is worth stating plainly because it runs the right way round: the data did not go looking for a service to justify. The scarcest skill in the set happens to sit in the workstream with the least tolerance for delay, and the mechanism for staffing it across a border already exists.

Conclusion: For a Fusion programme in the Gulf: staff finance locally, and treat payroll as a corridor role from day one rather than discovering it late.

Salesforce: why a Riyadh programme is a mobilisation planning problem

The local Salesforce pool in Saudi Arabia is small enough to count. That is not a reason to avoid the market; it is a reason to decide how the programme will be staffed at business case rather than at week six.

The extract finds 65 Salesforce professionals in the whole of Saudi Arabia, against 4,326 in the UK. Inside that 65: Marketing Cloud 18, Service Cloud 6, data migration 2, release engineering 1, and Commerce Cloud 0.

This is a supply measurement and nothing more. It says how many professionals declare the skill; it says nothing about how many roles are open, how long a hire takes, or how capable the market is. Saudi Arabia is building enterprise capability quickly and the picture will not look like this indefinitely.

What it does say is that a Salesforce programme in Riyadh cannot assume a local starting point, and a commerce programme has none at all. That is arithmetic, and arithmetic is planning information.

Two programmes with identical scope and identical budget can have very different outcomes here, and the difference is decided before either one starts.

The first plans for a local team, mobilises, discovers in week six that the shortlist is not going to materialise, and then does the corridor sourcing anyway — late, at pace, against a date that has already been communicated. The second writes the corridor into the business case, budgets for mobilisation and entity costs, and starts sourcing before the programme does.

Nothing in the data distinguishes those two programmes. The only difference is whether the plan admitted the arithmetic while it was still cheap to admit it.

Yallo operates across the UK, the Middle East and India because programmes routinely need specialists who are not in the country the programme is in. That is not a claim about any market's capability. It is a statement about how enterprise platform skills are distributed at a given moment, and this measurement is one snapshot of that distribution.

The practical form it takes is unglamorous: name the roles that will have to come from outside, price mobilisation and employment properly, and start those searches first because they have the longest lead time.

Conclusion: A Salesforce programme in Riyadh is a corridor programme by arithmetic. The planning question is not whether specialists come from outside, it is whether the plan says so at business case or discovers it at week six.

Cloud and DevOps: a pool broad enough to be a poor filter

This is the one layer where declared skills are genuinely broad. Breadth moves the entire burden from sourcing to screening, and it hides a regional pattern that runs against the received assumption.

Azure DevOps Services is declared by 48.7% of the 25,262 cloud and DevOps professionals measured, and AWS by 45.5%. Those are the two highest shares in this family.

They are not the highest in the whole extract — Salesforce Marketing Cloud+AMPscript reaches 60.3% — and the distinction matters, because comparing shares across families compares the width of two different title nets rather than two markets.

A skill declared by nearly half a pool has stopped discriminating. A search on it returns a list that is large and undifferentiated, which feels like abundance and behaves like noise. The constraint at this layer is not finding candidates; it is telling them apart, and that is screening work rather than sourcing work.

This is the clearest place in the corpus where the gap between a declared skill and a demonstrated one is the entire problem. Everything on LinkedIn is self-declared. A profile listing a platform is not a person who has run it in production under load.

The received view is that the Gulf is Azure-first. On this measure Google Cloud is declared by 19.1% of the Saudi pool against 10.4% of the UK pool — 6.9 percentage points higher across the Gulf as a whole.

One measurement does not overturn a market assumption, and this one is self-declared. It is enough to say that a Saudi programme choosing its platform on an assumption about local skills availability should check the assumption rather than inherit it.

The source net for this family is cloud and DevOps titles together, and that spans two of our desks rather than one. This piece is keyed to cloud and infrastructure and cross-links platform engineering and DevOps, because the net is what was measured and splitting it after the fact to produce a tidier mapping would be inventing precision the data does not have.

Conclusion: When almost half a pool declares the same skill, the declaration stops discriminating and the work moves to screening. Sourcing is not the constraint at this layer.

AI and data: the distance between a title and a tool

The market has data people. It has far fewer data platform people, and fewer still who have touched a named AI platform. At this end of the market, hiring stops being a funnel.

128,813 data professionals across the three markets. Databricks is declared by 3.6% of them and Snowflake by 3.2%.

The market has data people. It does not, on this measure, have data platform people in anything like the same quantity. That is the single clearest expression in the corpus of the problem this business exists to solve: a title is abundant and a capability is not, and only one of the two is visible in a search.

Across a pool of 22,853 AI and machine learning professionals: Amazon Bedrock 178, Vertex AI 206, Azure AI Foundry 189. Each is under one per cent of the pool.

Those three counts sum to 573 declarations, and that is the number to be careful with. A professional can declare more than one of these platforms, so the count of distinct individuals is lower than 573 and the extract cannot say how much lower. Even taken as an upper bound it is a small number for three national markets.

This is what replaces the adjective. Claims about AI talent nobody else can find are ordinarily unfalsifiable; this is a measurement, with a date on it and a stated method, and it can be checked.

The practical consequence is a change of kind rather than degree. A funnel assumes enough candidates at the top that filtering produces a shortlist. At these quantities there is no funnel to filter. The work becomes identifying named individuals, understanding what they have actually built, and approaching them properly, which is a different activity with different economics.

Conclusion: For AI platform work the pool is small enough to be a named-individual exercise rather than a funnel, and the figure replaces the adjective in claims about hard-to-find AI talent.

What this does not tell you

It counts professionals who exist, not roles that are open, so it supports no statement about how long a hire takes or which role is hardest to fill. Those need a firm's own placement data, and ours is not published here.

It carries nothing about compensation. There is no pay data in the source at all, so any rate or salary figure attached to this analysis did not come from it.

It is a snapshot dated 2 August 2026. Markets that are building capability quickly will not match it for long, and Saudi Arabia in particular should be expected to move.

YALLO TALENT · YALLO.CO · FIGURES FROM A LINKEDIN TALENT INSIGHTS EXTRACT READ ON 2 AUGUST 2026. SKILLS ARE SELF-DECLARED. COUNTS WITHIN A FAMILY OVERLAP AND ARE NOT PARTS OF A WHOLE. THIS IS SUPPLY-SIDE DATA AND SUPPORTS NO CLAIM ABOUT HOW LONG A ROLE TAKES TO FILL.

Telco & Media

Education & Universities

• COMPANY

About Yallo

Why Yallo

Leadership

Insights

Case studies

Jobs

• GROUP COMPANIES

Yallo Talent

saasinator AI

Privacy · Terms · Cookies